

Working With Panel Data

Introduction

Panel data refers to data with multiple observations per unit. In education settings panel data is almost more common than not, with many studies involving cases that have been observed over time.

For all of the models below, I'll use the following notation:

y_{it} is the dependent variable for unit i ($i = 1 \dots n$) in time period t ($t = 1 \dots t$).

x_{it} is an independent variable for unit i at time t .

β is a coefficient on the variable x

ϵ_{it} is an error term

The terminology around panel data can be confusing, because economists and education experts discuss the same things using different names. Here's some terminology:

Panel data: when used by economists, this typically refers to a dataset where there are many more units than observations over time.

Cross-sectional time-series data: this refers to data where there are much longer time series, and fewer data points.

Hierarchical or "grouped" data: this refers to data where the observations are naturally grouped, e.g. students in classrooms, classrooms in schools. This type of data can also include multiple observations over time.

Fixed effects: when used by economists, this refers to models where the group mean is controlled for, either by subtracting it from the dependent variable or by individually controlling for each group effect via dummy variables. Also known as LSDV: least squares dummy variables. When HLM people say fixed effects, they're referring to coefficients that don't vary across groups. This is also known as a "no pooling" model.

Random effects: when used by economists, this refers to a model that allows one or more coefficients to have its own distribution with an error term. A random effects model is functionally equivalent to a Hierarchical Linear Model, although HLM imposes additional assumptions.

Panel Data Model

$$y_{it} = \alpha_i + \beta_1 x_{1it} + \dots \beta_k x_{kit} + (\mu_i + \epsilon_{it}) \quad (1)$$

Fixed effect: α_{it}

Random effect: μ_{it}

Hierarchical Linear Model

$$y_{it} = \alpha_i + \beta_1 x_{1it} + \dots \beta_k x_{kit} + \epsilon_{it}$$

Where:

$$\alpha_i = \gamma_0 + \gamma_1 z_i + \mu_i \quad (2)$$

Fixed effects: $\beta_1 \dots \beta_k$

Random effect: α_i

Describing Panel Data

The data we'll be using come from my dissertation, which prediction appropriations, tuition and financial aid at the state level using various characteristics of the political and higher education system. The data are a balanced panel of 49 states (excluding Alaska) over 16 years, 1984-1999.

To get Stata to recognize this as panel data, we need to use the xtset command.

```
. /* Set up data as panel data */  
. xtset state year, yearly  
      panel variable:  state (strongly balanced)  
      time variable:  year, 1984 to 1999  
      delta:  1 year
```

I tend to use two basic methods for describing panel data. First, I like to do line graphs for all of the continuous variables, which give you a very clear sense of variation across units and any time trends. It's also a good way to find data problems:

```
. xtline approps_i  
  
. xtline pub4tuit_i
```

The other graph I like to use is a boxplot for the variable by state. This gives an excellent sense of variability both across and within units.

```
.  
. #delimit ;  
delimiter now ;  
. graph hbox pub4tuit_i,  
> over(state, sort(1) descending label(labsize(tiny) ))  
>  
> ;
```




Figure 1: Trend in Appropriations Per Student, by State

```
. #delimit ;
delimiter now ;
. graph hbox approps,
> over(state, sort(1) descending label(labsize(tiny) ))
> ;
```

When reporting descriptives for a panel dataset, don't just give the grand mean. Provide averages and standard deviations for a subset of time periods, along with graphics similar to the above.

Ordinary Least Squares

The OLS estimate for panel data is:

$$y_{it} = \alpha + \beta x_{it} + \epsilon_{it}$$

In Stata:

```
.
```




Figure 2: Trend in Public Four-Year Tuition, by State

```
. local y approps_i
.
. local controls perc1824 incpcp_i percpriv taxcpc_i legcomp_i i.board
.
. reg `y' legideo `controls'
```

| Source | SS | df | MS | Number of obs = | 784 |
|----------|------------|-----|------------|-----------------|--------|
| Model | 830964034 | 10 | 83096403.4 | F(10, 773) = | 111.00 |
| Residual | 578657976 | 773 | 748587.29 | Prob > F = | 0.0000 |
| Total | 1.4096e+09 | 783 | 1800283.54 | R-squared = | 0.5895 |
| | | | | Adj R-squared = | 0.5842 |
| | | | | Root MSE = | 865.21 |

| approps_i | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] |
|-----------|-----------|-----------|--------|-------|----------------------|
| legideo | 1.954075 | 1.373991 | 1.42 | 0.155 | -.7431209 4.651272 |
| perc1824 | 267.1039 | 29.04555 | 9.20 | 0.000 | 210.0865 324.1214 |
| incpcp_i | -12.65837 | 12.94071 | -0.98 | 0.328 | -38.06147 12.74473 |
| percpriv | -57.78148 | 2.810894 | -20.56 | 0.000 | -63.29937 -52.26359 |
| taxcpc_i | 1.939732 | .1145756 | 16.93 | 0.000 | 1.714816 2.164649 |
| legcomp_i | -.0008065 | .0020159 | -0.40 | 0.689 | -.0047638 .0031508 |
| board | | | | | |
| 2 | 110.3047 | 100.8765 | 1.09 | 0.275 | -87.71972 308.3291 |
| 3 | -28.19471 | 94.82565 | -0.30 | 0.766 | -214.341 157.9516 |




Figure 3: Variation in Appropriations Per Student, by State

| | | | | | | | |
|-------|--|-----------|----------|--------|-------|-----------|-----------|
| 4 | | -29.0085 | 87.02584 | -0.33 | 0.739 | -199.8435 | 141.8265 |
| 5 | | -1538.795 | 143.7325 | -10.71 | 0.000 | -1820.947 | -1256.643 |
| _cons | | 944.5651 | 466.3937 | 2.03 | 0.043 | 29.01674 | 1860.113 |

The problem with the OLS model is both that it may be inconsistent and that it may induce huge problems with heteroscedasticity. If you're not sure if you there's a problem, try graphing the residuals like so:

```
. predict e, resid
.
. graph box e, over(state, sort(1) descending label(labsize(tiny))) /*Horrible*/
```

In our case, there are massive problems with the error terms by state. It's not so bad by year. Even so, we will have a correlation with the independent variables and the error term because we're leaving out a variable that is known to impact the dependent variable: the group that each unit is in.




Figure 4: Variation in Public Four-Year Tuition, by State

Fixed Effects Models

The fixed effects model with group specific intercepts is:

$$y_{it} = \alpha_i + \beta x_{it} + \epsilon_{it}$$

A basic fixed effects model looking at the effect of a more liberal government on appropriations would be specified as:

```
. xi: xtreg `y' legideo `controls', fe
i.board      _Iboard_1-5      (naturally coded; _Iboard_1 omitted)
note: _Iboard_5 omitted because of collinearity

Fixed-effects (within) regression               Number of obs   =       784
Group variable: state                          Number of groups =       49

R-sq:  within = 0.2281                        Obs per group:  min =      16
        between = 0.0860                                avg  =     16.0
        overall  = 0.1015                                max  =      16

corr(u_i, Xb)  = -0.2562                      F(9,726)        =     23.83
                                                Prob > F         =     0.0000
```




Figure 5: Residuals by State

| approps_i | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|--|--------------|--------------------------------------|-------|---------------------|----------------------|-----------|
| legideo | 3.508206 | 1.188178 | 2.95 | 0.003 | 1.175531 | 5.840881 |
| perc1824 | 269.2799 | 24.15645 | 11.15 | 0.000 | 221.8551 | 316.7048 |
| incpcp_i | 12.85631 | 20.06298 | 0.64 | 0.522 | -26.53207 | 52.24468 |
| percpriv | -3.949699 | 10.39503 | -0.38 | 0.704 | -24.35761 | 16.45821 |
| taxcpc_i | 1.436178 | .1566408 | 9.17 | 0.000 | 1.128655 | 1.743701 |
| legcomp_i | .0022197 | .001997 | 1.11 | 0.267 | -.0017009 | .0061402 |
| _Iboard_2 | -41.50431 | 108.1897 | -0.38 | 0.701 | -253.9063 | 170.8977 |
| _Iboard_3 | -597.6449 | 204.1981 | -2.93 | 0.004 | -998.5342 | -196.7557 |
| _Iboard_4 | -942.1278 | 183.3558 | -5.14 | 0.000 | -1302.099 | -582.1569 |
| _Iboard_5 | (omitted) | | | | | |
| _cons | -38.59943 | 659.1182 | -0.06 | 0.953 | -1332.605 | 1255.406 |
| sigma_u | 1232.7514 | | | | | |
| sigma_e | 492.51715 | | | | | |
| rho | .86235025 | (fraction of variance due to u_i) | | | | |
| ----- | | | | | | |
| F test that all u_i=0: | | F(48, 726) = | 34.57 | | Prob > F | |
| = 0.0000 | | | | | | |
| ----- | | | | | | |
| . xi: reg `y' legideo `controls' i.state | | | | | | |
| i.board | _Iboard_1-5 | (naturally coded; _Iboard_1 omitted) | | | | |
| i.state | _Istate_2-50 | (naturally coded; _Istate_2 omitted) | | | | |
| note: _Istate_22 omitted because of collinearity | | | | | | |
| Source | SS | df | MS | Number of obs = 784 | | |




Figure 6: Residuals by Year

| | | | | | | |
|-------------|------------|-----------|------------|------------------------|-------|----------------------|
| -----+----- | | | | F(57, 726) = 89.21 | | |
| Model | 1.2335e+09 | 57 | 21640594.9 | Prob > F = 0.0000 | | |
| Residual | 176108102 | 726 | 242573.144 | R-squared = 0.8751 | | |
| -----+----- | | | | Adj R-squared = 0.8653 | | |
| Total | 1.4096e+09 | 783 | 1800283.54 | Root MSE = 492.52 | | |
| -----+----- | | | | | | |
| approps_i | | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] |
| -----+----- | | | | | | |
| legideo | | 3.508206 | 1.188178 | 2.95 | 0.003 | 1.175531 5.840881 |
| perc1824 | | 269.2799 | 24.15645 | 11.15 | 0.000 | 221.8551 316.7048 |
| incpcp_i | | 12.85631 | 20.06298 | 0.64 | 0.522 | -26.53207 52.24468 |
| percpriv | | -3.949699 | 10.39503 | -0.38 | 0.704 | -24.35761 16.45821 |
| taxpc_i | | 1.436178 | .1566408 | 9.17 | 0.000 | 1.128655 1.743701 |
| legcomp_i | | .0022197 | .001997 | 1.11 | 0.267 | -.0017009 .0061402 |
| _lboard_2 | | -41.50431 | 108.1897 | -0.38 | 0.701 | -253.9063 170.8977 |
| _lboard_3 | | -597.6449 | 204.1981 | -2.93 | 0.004 | -998.5342 -196.7557 |
| _lboard_4 | | -942.1278 | 183.3558 | -5.14 | 0.000 | -1302.099 -582.1569 |
| _lboard_5 | | -1876.622 | 218.9806 | -8.57 | 0.000 | -2306.533 -1446.711 |
| _lstate_3 | | 276.3964 | 184.0882 | 1.50 | 0.134 | -85.01223 637.8051 |
| _lstate_4 | | -933.8803 | 259.3431 | -3.60 | 0.000 | -1443.032 |
| -----+----- | | | | | | |
| -424.7282 | | | | | | |
| -----+----- | | | | | | |
| | | | | | | |

This includes both the standard xtreg command and a reg command, with xi specified

to control for state level effects. The coefficients are the same. The interpretation of a fixed effects model always refers only to within-unit changes in both the independent and dependent variables.

Without correcting for time in the above model, we could introduce serially correlated error terms.

Fixed Effects for Time

In addition to specifying fixed effects for groups, the simplest approach to handling time is to specify fixed effects for time, with T-1 variables for time included in the model, with a new set of coefficients γ_t .

$$y_{it} = \alpha_i + \beta x_{it} + \gamma_t + \epsilon_{it}$$

To estimate the above in stata, we would need to use the `i` function, which transforms variables into a categorical variable. The following syntax gives fixed effects for time, with time as a categorical variable:

```
. /* Fixed Effects for Units and Time (state and year) */
.
. xtreg `y' legideo `controls' i.year , fe
note: 5.board omitted because of collinearity
```

| | | | |
|-----------------------------------|----------------------|---|--------|
| Fixed-effects (within) regression | Number of obs | = | 784 |
| Group variable: state | Number of groups | = | 49 |
| R-sq: within = 0.3942 | Obs per group: min = | | 16 |
| between = 0.0321 | avg = | | 16.0 |
| overall = 0.0576 | max = | | 16 |
| | F(24,711) | = | 19.27 |
| corr(u_i, Xb) = -0.4822 | Prob > F | = | 0.0000 |

| approps_i | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] |
|-----------|-----------|-----------|-------|-------|----------------------|
| legideo | 1.145978 | 1.140491 | 1.00 | 0.315 | -1.093154 3.385111 |
| perc1824 | 44.91926 | 30.75181 | 1.46 | 0.145 | -15.45595 105.2945 |
| incpcp_i | 139.2413 | 28.45662 | 4.89 | 0.000 | 83.37225 195.1104 |
| percpriv | -3.777036 | 10.16795 | -0.37 | 0.710 | -23.73984 16.18577 |
| taxcpc_i | 1.501035 | .1425019 | 10.53 | 0.000 | 1.22126 1.78081 |
| legcomp_i | .0016732 | .0018213 | 0.92 | 0.359 | -.0019025 .005249 |
| board | | | | | |
| cbweak | -24.88159 | 97.00964 | -0.26 | 0.798 | -215.3412 165.578 |
| gball | -454.4452 | 183.9582 | -2.47 | 0.014 | -815.6115 -93.27897 |
| gbfour | -711.7997 | 165.6897 | -4.30 | 0.000 | -1037.099 -386.5002 |
| plan | 0 | (omitted) | | | |
| year | | | | | |
| 1985 | 220.3299 | 90.99408 | 2.42 | 0.016 | 41.68063 398.9791 |
| 1986 | 113.4229 | 97.881 | 1.16 | 0.247 | -78.74752 305.5932 |
| 1987 | -19.38342 | 105.0826 | -0.18 | 0.854 | -225.6928 186.926 |
| 1988 | -67.57019 | 113.1309 | -0.60 | 0.551 | -289.6807 154.5403 |
| 1989 | -274.3213 | 122.1241 | -2.25 | 0.025 | -514.0883 -34.55431 |
| 1990 | -399.1526 | 125.7373 | -3.17 | 0.002 | -646.0135 -152.2917 |
| 1991 | -657.3481 | 125.0003 | -5.26 | 0.000 | -902.7619 -411.9343 |

| | | | | | | | |
|---|--|-----------|-----------------------------------|-------|-------|-------------------|-----------|
| 1992 | | -678.0808 | 134.1735 | -5.05 | 0.000 | -941.5044 | -414.6571 |
| 1993 | | -936.106 | 136.6561 | -6.85 | 0.000 | -1204.404 | -667.8083 |
| 1994 | | -968.5213 | 145.4102 | -6.66 | 0.000 | -1254.006 | -683.0365 |
| 1995 | | -1031.559 | 153.3461 | -6.73 | 0.000 | -1332.624 | -730.4935 |
| 1996 | | -1044.886 | 162.4511 | -6.43 | 0.000 | -1363.827 | -725.9445 |
| 1997 | | -1058.236 | 171.5086 | -6.17 | 0.000 | -1394.96 | -721.5123 |
| 1998 | | -1197.384 | 189.6214 | -6.31 | 0.000 | -1569.669 | -825.0989 |
| 1999 | | -1194.228 | 195.1562 | -6.12 | 0.000 | -1577.379 | -811.0763 |
| | | | | | | | |
| _cons | | -163.0829 | 776.2306 | -0.21 | 0.834 | -1687.061 | 1360.895 |
| ----- | | | | | | | |
| sigma_u | | 1421.003 | | | | | |
| sigma_e | | 440.90254 | | | | | |
| rho | | .91218326 | (fraction of variance due to u_i) | | | | |
| ----- | | | | | | | |
| F test that all u_i=0: F(48, 711) = 52.77 | | | | | | Prob > F = 0.0000 | |

The interpretation of this would be as usual for a categorical variable: each coefficient for time represents a contrast to a base time period (stata will choose the first one). Having done this however, concerns about serial correlation should be adequately addressed.

This can be observed by looking at a boxplot of errors by time period:

```
predict res,e graph box res, over(year)
```

Fixed effects for time are not symmetric with fixed effects for groups in this model. To adjust for this, we can regress

$$y_{*it} = y_{it} - \bar{y}_i - \bar{y}_t + \bar{y}$$

on the independent variable x, specified as:

$$x_{*it} = x_{it} - \bar{x}_i - \bar{x}_t + \bar{x}$$

Serially Correlated Errors

Fixed effects for time is an appropriate approach in many cases, however it is very inefficient: if time itself is not of interest, you will have $T - 1$ nuisance parameters along with $n - 1$ group estimates in the case of a fixed effects approach.

When estimating models for panel data, corrections for autocorrelation are much the same as in a single sample. First, assume that there is no cross-sectional autocorrelation:

$$\text{Corr}[\epsilon_{it}, \epsilon_{js}] = 0, \text{ if } i \neq j$$

In the presence of within-unit autocorrelation, the observed error ϵ_{it} consists of two parts: the error term in the previous year multiplied by a coefficient ρ and the overall error term μ_{it} .

$$\epsilon_{it} = \rho_i \epsilon_{it-1} + \mu_{it}$$

The variance of these group-specific error terms is therefore:

$$\text{Var}[e_{it}] = \sigma_i^2 = \frac{\sigma_\mu^2 i}{1 - \rho_i^2}$$

To account for this, we need to calculate a correlation coefficient ρ for each group. A group specific estimate r_i for ρ is:

$$r_i = \frac{\sum_{t=2}^T e_{it} e_{i,t-1}}{\sum_{t=1}^T e_{it}^2}$$

Most programs, including STATA, calculate a single value, which is the average of all group specific correlation coefficients. This value is then used to transform the data to eliminate the autocorrelation. For instance for y_{it} , the transformation is:

$$y_{i1}, y_{i2}, \dots, y_{iT} = \sqrt{1 - r^2} y_{i1}, y_{i2} - r_i y_{i1}, y_{i3} - r_i y_{i2}, y_{iT} - r_i y_{i,T-1}$$

To estimate a fixed effects model in STATA, use the `xtregar` command. In our running example, this can be estimated via:

```
. xi: xtregar `y' legideo `controls', fe rhotype (tsc) twostep lbi
i.board      _Iboard_1-5      (naturally coded; _Iboard_1 omitted)
note: _Iboard_5 dropped because of collinearity
```

| | | | |
|--|----------------------|---|--------|
| FE (within) regression with AR(1) disturbances | Number of obs | = | 735 |
| Group variable: state | Number of groups | = | 49 |
| R-sq: within = 0.1475 | Obs per group: min = | | 15 |
| between = 0.0584 | avg = | | 15.0 |
| overall = 0.0235 | max = | | 15 |
| | F(9,677) | = | 13.02 |
| corr(u_i, Xb) = -0.6379 | Prob > F | = | 0.0000 |

| approps_i | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] |
|-----------|-----------|---------------------------------------|-------|-------|----------------------|
| legideo | 2.203477 | 1.328746 | 1.66 | 0.098 | -.4054819 4.812437 |
| perc1824 | 308.6154 | 33.12758 | 9.32 | 0.000 | 243.5702 373.6605 |
| incpcp_i | 16.72646 | 23.96317 | 0.70 | 0.485 | -30.32461 63.77753 |
| percpriv | 39.45439 | 12.99704 | 3.04 | 0.002 | 13.93505 64.97374 |
| taxcpc_i | .9035681 | .1776652 | 5.09 | 0.000 | .5547272 1.252409 |
| legcomp_i | .0015629 | .0019213 | 0.81 | 0.416 | -.0022095 .0053352 |
| _Iboard_2 | -88.92507 | 136.4487 | -0.65 | 0.515 | -356.8385 178.9884 |
| _Iboard_3 | -444.8078 | 257.5141 | -1.73 | 0.085 | -950.4301 60.8145 |
| _Iboard_4 | -797.8778 | 234.5891 | -3.40 | 0.001 | -1258.488 -337.268 |
| _Iboard_5 | (omitted) | | | | |
| _cons | -630.4428 | 486.6708 | -1.30 | 0.196 | -1586.008 325.1228 |
| rho_ar | .38558457 | | | | |
| sigma_u | 1626.9651 | | | | |
| sigma_e | 422.42905 | | | | |
| rho_fov | .93684349 | (fraction of variance because of u_i) | | | |

```
F test that all u_i=0:      F(48,677) =      22.21      Prob > F = 0.0000
modified Bhargava et al. Durbin-Watson = 1.0483739
Baltagi-Wu LBI = 1.2288309
```

However, the transformation of the data in the above is done via the Cochrane-Orcutt, not Prais-Winsten transformation. Cochrane-Orcutt throws out the first unit in each time series,

which can be a lot of data in a panel data setting. Another option is to use xtpcse, with correlation set to AR(1). This also incorporates some other assumptions, which can be turned off by specifying the “independent” option. In particular, this allows for unit-specific autocorrelation, which is generally a better assumption.

```
. // Unit-specific ar(1) process
. xtpcse `y' legideo `controls' i.state, correlation (psar1) independent
note: 46.state omitted because of collinearity
(note: estimates of rho outside [-1,1] bounded to be in the range [-1,1])

Prais-Winsten regression, independent panels corrected standard errors

Group variable:      state                Number of obs   =       784
Time variable:      year                Number of groups  =       49
Panels:             independent (balanced)  Obs per group:
Autocorrelation:    panel-specific AR(1)      min =       16
                                                avg  =       16
                                                max  =       16

Estimated covariances      =       1      R-squared      =       0.9520
Estimated autocorrelations =       49      Wald chi2(56)   =      2177.81
Estimated coefficients     =       57      Prob > chi2     =       0.0000
```

| approps_i | Indep-corrected | | | | | |
|-----------|-----------------|-----------|--------|-------|----------------------|-----------|
| | Coef. | Std. Err. | z | P> z | [95% Conf. Interval] | |
| legideo | 1.941346 | 1.113935 | 1.74 | 0.081 | -.2419273 | 4.124619 |
| perc1824 | 263.31 | 28.58896 | 9.21 | 0.000 | 207.2767 | 319.3434 |
| incpcp_i | 124.4886 | 21.42011 | 5.81 | 0.000 | 82.50597 | 166.4712 |
| percpriv | 1.636201 | 10.40599 | 0.16 | 0.875 | -18.75916 | 22.03156 |
| taxcpc_i | .5548492 | .1564309 | 3.55 | 0.000 | .2482503 | .8614482 |
| legcomp_i | .0029899 | .0017338 | 1.72 | 0.085 | -.0004084 | .0063881 |
| board | | | | | | |
| cbweak | -103.2285 | 116.8174 | -0.88 | 0.377 | -332.1863 | 125.7294 |
| gball | -1015.286 | 183.3439 | -5.54 | 0.000 | -1374.634 | -655.9389 |
| gbfour | -809.2099 | 194.6451 | -4.16 | 0.000 | -1190.707 | -427.7125 |
| plan | -4170.79 | 401.9689 | -10.38 | 0.000 | -4958.635 | -3382.945 |

Random Effects

In the random effects model, the group effect is assumed to have a distribution and an error term. You’ll get a LOT more on this in Regression II, so today I’ll just introduce it to you and show you how to run the Hausman test. In practice, a random effect model is rarely appropriate unless the groups are defined as part of the sampling procedure.

First Differenced Model

First differencing is just what it sounds like: subtracting the previous time period’s observation from the current one:

$$\Delta y_{it} = y_{it} - y_{it-1}$$

A first differenced model can be used with panel data, although the interpretation of coefficients goes from change in x to change-in-change in x .

$$\Delta y_{it} = \beta_0 + \beta_1 \Delta x_{it} + \epsilon$$